

Ensemble methods

Victor Kitov

v.v.kitov@yandex.ru

Yandex School of Data Analysis



Ensemble learning

Definition 1

Ensemble learning - using multiple machine learning methods for a given problem and combining their outputs to obtain final result.

Synonyms: committee-based learning, multiple classifier systems.

Motivation of ensembles

- Benefits for prediction:
 - increased accuracy
 - increased robustness.
- Justification: some predictors are compensating the errors of other predictors
- When to use:
 - existing model hypothesis space is too narrow to explain the true one (high model bias)
 - avoid local optima of optimization methods (high model variance)
 - too small dataset to figure out concretely the exact model hypothesis
- Frequently the task itself promotes usage of ensembles (such as computer security):
 - multiple sources of diverse information
 - different abstraction levels need to be united

Table of Contents

1 Motivation

- Motivation for classification
- Motivation for regression

2 Popular ensemble methods

Motivation

Motivation for classification

1

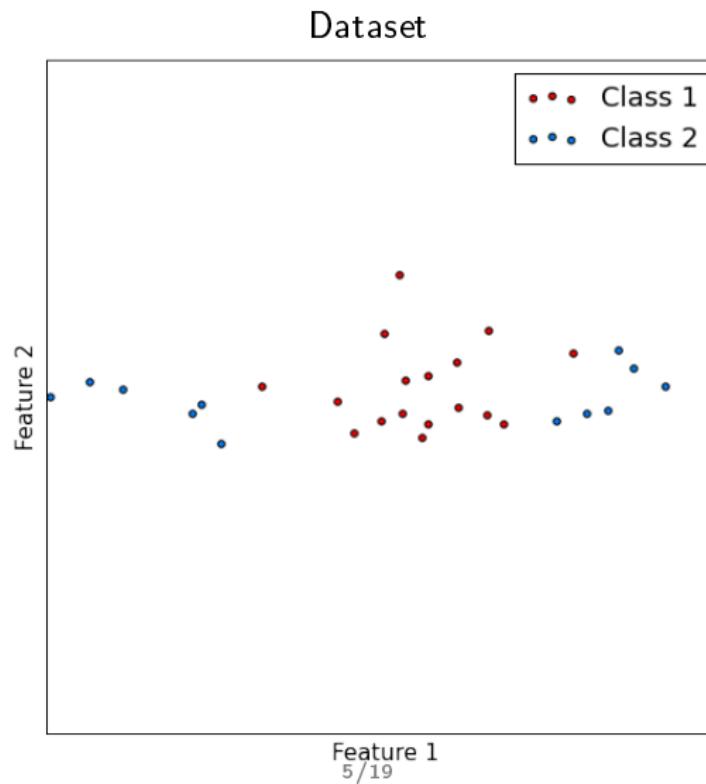
Motivation

- Motivation for classification
- Motivation for regression

Motivation

Motivation for classification

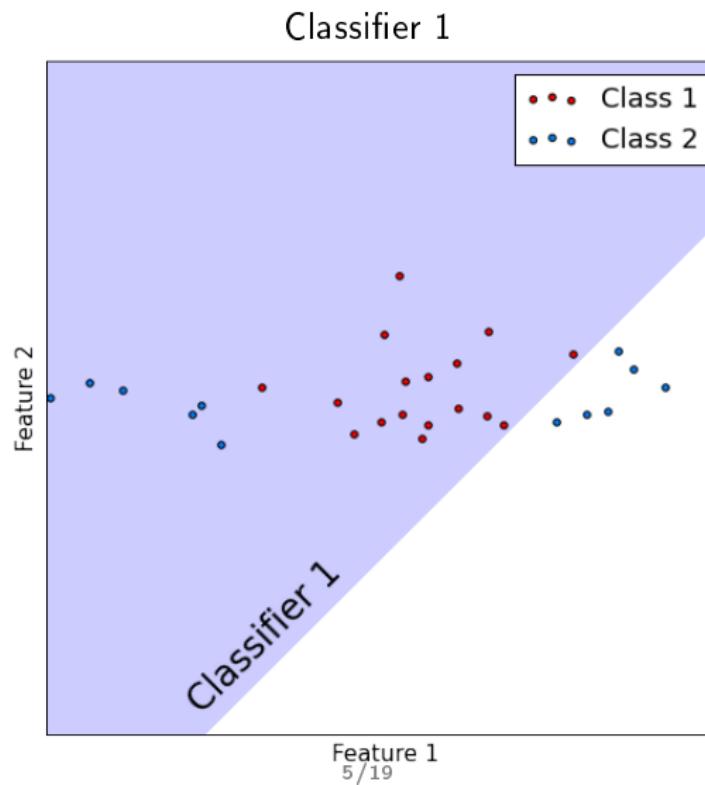
Motivation for classification



Motivation

Motivation for classification

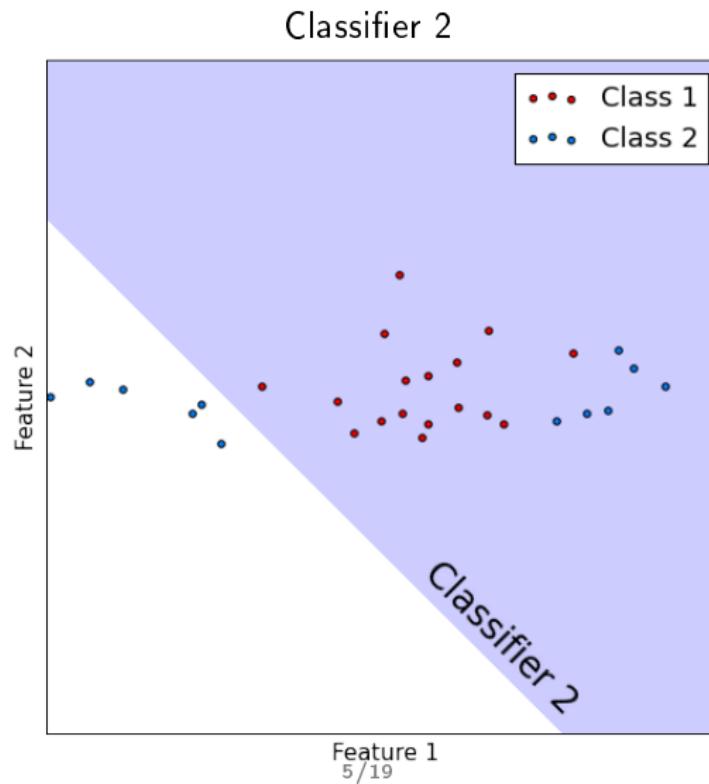
Motivation for classification



Motivation

Motivation for classification

Motivation for classification

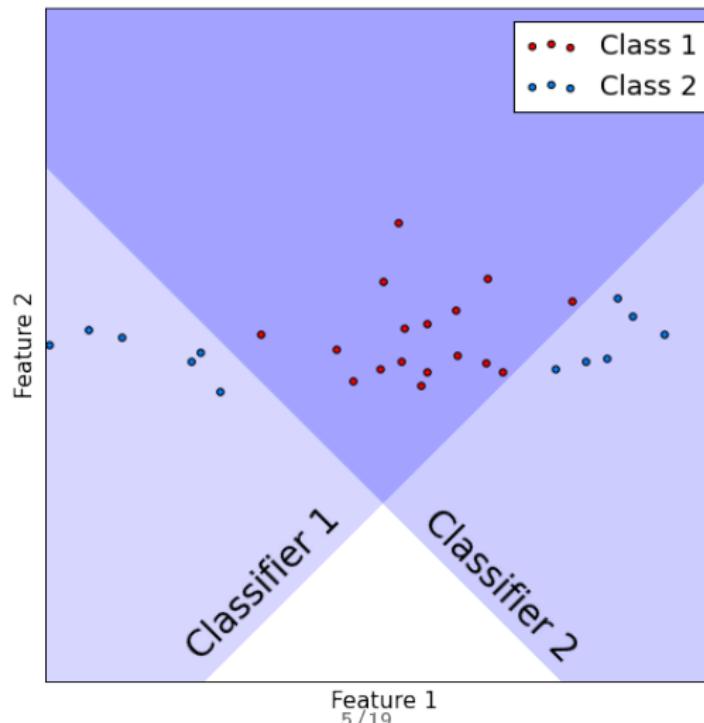


Motivation

Motivation for classification

Motivation for classification

Classifier 1 and classifier 2 combined using AND rule



Motivation

Motivation for regression

1 Motivation

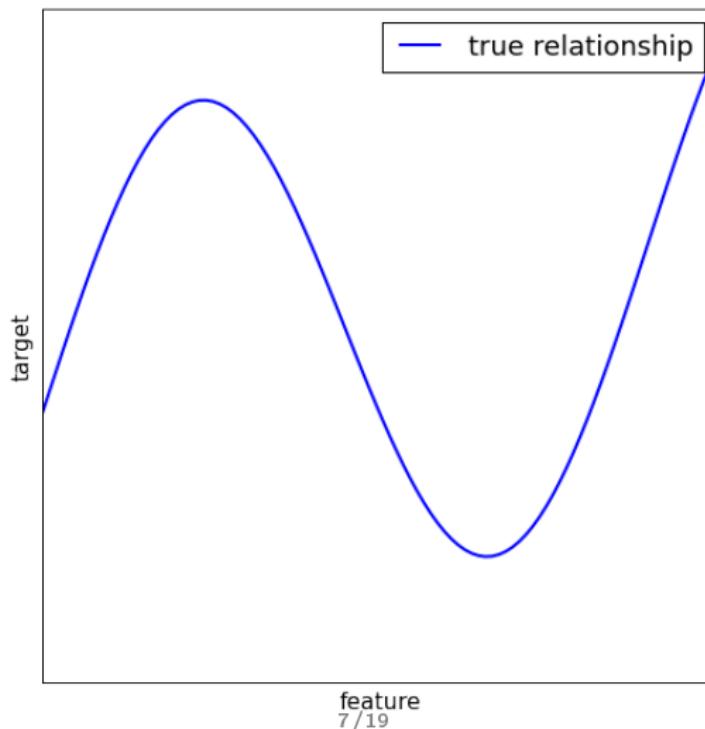
- Motivation for classification
- Motivation for regression

Motivation

Motivation for regression

Motivation for regression

Dataset

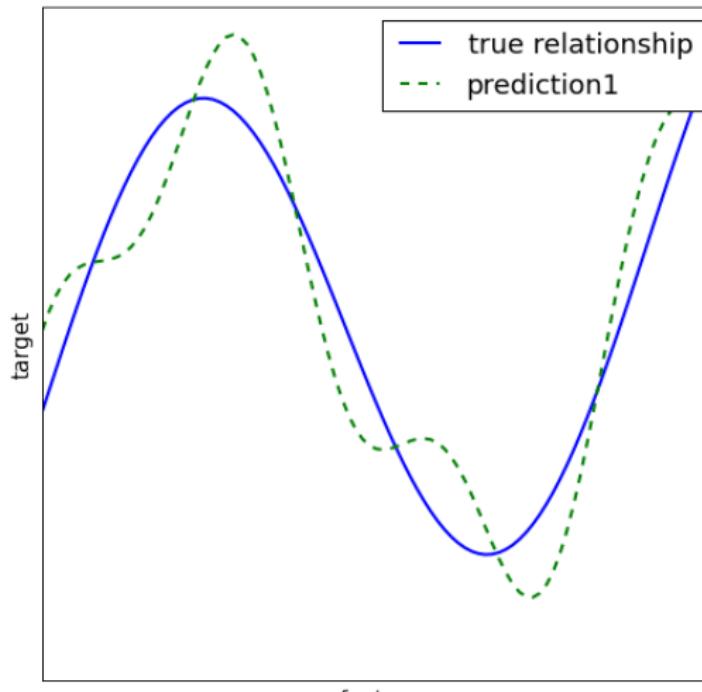


Motivation

Motivation for regression

Motivation for regression

Regression 1

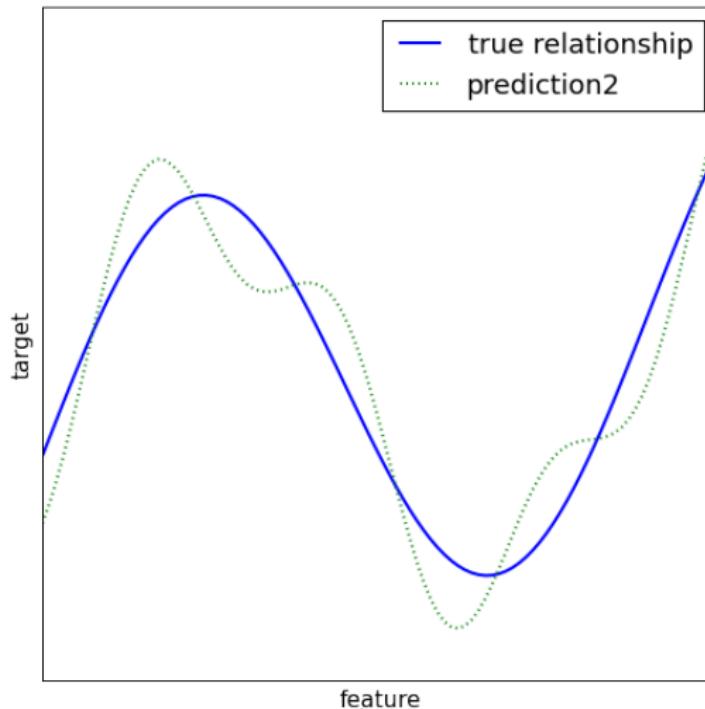


Motivation

Motivation for regression

Motivation for regression

Regression 2



Motivation

Motivation for regression

Motivation for regression

Regression 1 and regression 2 combined using averaging

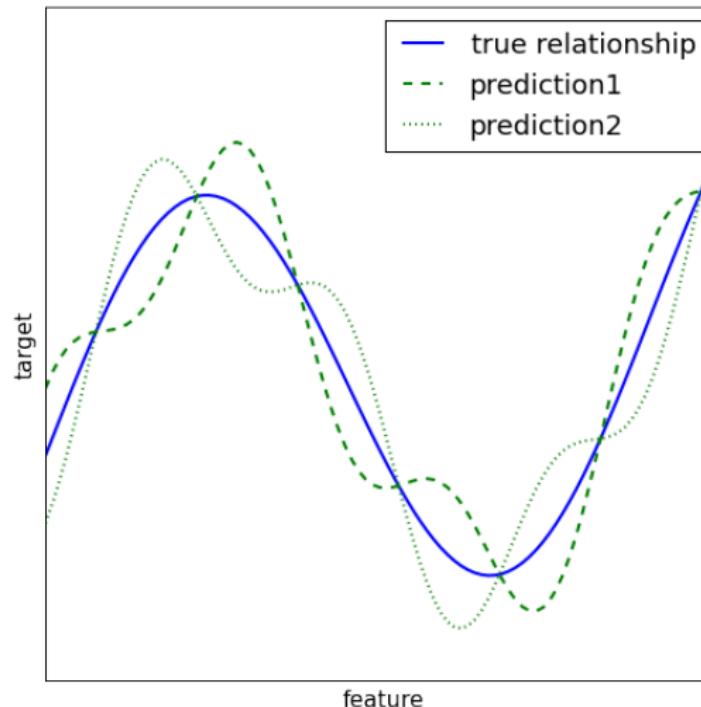


Table of Contents

1 Motivation

2 Popular ensemble methods

- Bagging and random forest
- Fixed integration schemes for classification

2 Popular ensemble methods

- Bagging and random forest
- Fixed integration schemes for classification

Bagging & random subspaces

- Bagging
 - random selection of samples (with replacement)
 - *what is the probability that observation will not belong to bootstrap sample?*
 - *what is the limit of this probability with $N \rightarrow \infty$?*
- Random subspace method:
 - random selection of features (without replacement)
- We can apply both methods jointly

Random forests

Input: training dataset $TDS = \{(x_i, y_i), 1 = 1, 2, \dots, N\}$; the number of trees B and the size of feature subsets m .

- ① for $b = 1, 2, \dots, B$:
 - ① generate random training dataset TDS^b of size N by sampling (x_i, y_i) pairs from TDS with replacement.
 - ② build a tree using TDS^b training dataset with feature selection for each node from random subset of features of size m (generated individually for each node) without replacement.
- ② Evaluate the quality by assigning output to $x_i, i = 1, 2, \dots, n$ using majority vote (classification) or averaging (regression) among trees with $b \in \{b : (x_i, y_i) \notin T^b\}$

Output: B trees. Classification is done using majority vote and regression using averaging of B outputs.

Comments

- Random forests use random selection on both samples and features
- Left out samples are used for evaluation of model performance.
- Less interpretable than individual trees
- +: Parallel implementation
- -: different trees are not targeted to correct mistakes of each other
- Extra-Random trees: more bias and less variance by random sampling (feature,value) pairs.

2 Popular ensemble methods

- Bagging and random forest
- Fixed integration schemes for classification

Fixed combiner at class level

Output of base learner k

Exact class: ω_1 or ω_2 .

Combiner predicts ω_1 if:

- all classifiers predict ω_1 (AND rule)
- at least one classifier predicts ω_1 (OR rule)
- at least k classifiers predict ω_1 (k-out-of-N)
- majority of classifiers predict ω_1 (majority vote)

Each classifier may be assigned a weight, based on its performance:

- weighted majority vote
- weighted k-out-of-N (based on score sum)

Fixed combiner - ranking level

Output of base learner k

Ranking of classes:

$$\omega_{k_1} \succeq \omega_{k_2} \succeq \dots \succeq \omega_{k_C}$$

Ranking is equivalent to scoring of each class (with incomparable scoring between classifiers).

Definition 2

Let $B_k(i)$ be the count of classes scored below ω_i by classifier k . **Borda count** $B(i)$ of class ω_i is the total number of classes scored below ω_i by all classifiers:

$$B(i) = \sum_{k=1}^K B_k(i)$$

Combiner predicts ω_i where $i = \arg \max_i B(i)$

Fixed combiner at class probability level

Output of base learner k

Vectors of class probabilities:

$$[p^k(\omega_1), p^k(\omega_2), \dots p^k(\omega_C)]$$

Combiner predicts ω_i if $i = \arg \max_i F(p^1(\omega_i), p^2(\omega_i), \dots p^K(\omega_i))$

- $F = \text{mean or median.}$

Finding constant weights

Weighted averaging combiner

$$f(x) = \sum_{k=1}^K w_k f_k(x)$$

Naive fitting

$$\hat{w} = \arg \min_w \sum_{i=1}^N \mathcal{L}(y_i, \sum_{k=1}^K w_k f_k(x_i))$$

will overfit. The mostly overfitted method will get the most weight.

Linear stacking

- Let training set $\{(x_i, y_i), i = 1, 2, \dots, N\}$ be split into M folds.
- Denote $fold(i)$ to be the fold, containing observation i
- Denote $f_k^{-fold(i)}$ be predictor k trained on all folds, except $fold(i)$.

Definition

Linear stacking (or stacked generalization) is weighted averaging combiner, where weights are found using

$$\hat{w} = \arg \min_w \sum_{i=1}^N \mathcal{L}(y_i, \sum_{k=1}^K w_k f_k^{-fold(i)}(x_i))$$

Generalized stacking

Definition

Generalized stacking is prediction

$$f(x) = A_\theta(f_1(x), f_2(x), \dots, f_K(x)),$$

where A is some general form predictor and θ is a vector of parameters, estimated by

$$\hat{\theta} = \arg \min_{\theta} \sum_{i=1}^N \mathcal{L} \left(y_i, A_\theta \left(f_1^{-fold(i)}(x), f_2^{-fold(i)}(x), \dots, f_K^{-fold(i)}(x) \right) \right)$$

- Stacking is the most general approach
- It is a winning strategy in most ML competitions.
- $f_i(x)$ may be:
 - class number (coded using one-hot encoding).
 - vector of class probabilities
 - any initial or generated feature