

# Boosting

Victor Kitov

v.v.kitov@yandex.ru

Yandex School of Data Analysis



## General idea

- Bagging, Random forest, extra random trees fit independent models.
  - can be done in parallel
  - sampling: objects - with replacement, features - without replacement.
  - base learners-complex
- Boosting fits models sequentially and they depend on each other.
  - each model is fitted to correct errors of the ensemble of previous models.
  - base learners-simple

# Linear ensembles

**Linear ensemble:**

$$F(x) = f_0(x) + c_1 h_1(x) + \dots + c_M h_M(x)$$

**Regression:**  $\hat{y}(x) = F(x)$

**Binary classification:**  $score(y|x) = F(x)$ ,  $\hat{y}(x) = \text{sign } F(x)$

- Notation:  $h_1(x), \dots, h_M(x)$  are called *base learners, weak learners, base models.*
- Too expensive to optimize  $f_0(x), h_1(x), \dots, h_M(x)$  and  $c_1, \dots, c_M$  jointly for large  $M$ .
- Idea: optimize  $f_0(x)$  and then each pair  $(h_m(x), c_m)$  greedily.

## Forward stagewise additive modeling (FSAM)

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$ , parametrized base learner  $h(x|\theta)$  and initial approximation  $f(x|\theta')$ , the number  $M$  of successive additive approximations.

# Forward stagewise additive modeling (FSAM)

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$ , parametrized base learner  $h(x|\theta)$  and initial approximation  $f(x|\theta')$ , the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x) = \arg \min_f \sum_{i=1}^N \mathcal{L}(f(x_i), y_i)$

# Forward stagewise additive modeling (FSAM)

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$ , parametrized base learner  $h(x|\theta)$  and initial approximation  $f(x|\theta')$ , the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x) = \arg \min_f \sum_{i=1}^N \mathcal{L}(f(x_i), y_i)$
- ② For  $m = 1, 2, \dots, M$ :
  - ① find next best classifier

$$(c_m, h_m) = \arg \min_{h,c} \sum_{i=1}^N \mathcal{L}(f_{m-1}(x_i) + ch(x_i), y_i)$$

# Forward stagewise additive modeling (FSAM)

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$ , parametrized base learner  $h(x|\theta)$  and initial approximation  $f(x|\theta')$ , the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x) = \arg \min_f \sum_{i=1}^N \mathcal{L}(f(x_i), y_i)$
- ② For  $m = 1, 2, \dots, M$ :

- ① find next best classifier

$$(c_m, h_m) = \arg \min_{h,c} \sum_{i=1}^N \mathcal{L}(f_{m-1}(x_i) + ch(x_i), y_i)$$

- ② set

$$f_m(x) = f_{m-1}(x) + c_m h_m(x)$$

**Output:** approximation function  
 $f_M(x) = f_0(x) + \sum_{m=1}^M c_m h_m(x)$

## Comments on FSAM

- Number of steps  $M$  should be determined by performance on validation set.
- Step 1 need not be solved accurately, since its mistakes are expected to be corrected by future base learners.
  - we can take  $f_0(x) = \arg \min_{\beta \in \mathbb{R}} \sum_{i=1}^N \mathcal{L}(\beta, y_i)$  or simply  $f_0(x) \equiv 0$ .
- By similar reasoning there is no need to solve 2.1 accurately
  - typically very simple base learners are used such as trees of depth=1,2,3.
- For some loss functions, such as  $\mathcal{L}(y, f(x)) = e^{-yf(x)}$  we can solve FSAM explicitly.
- For general loss functions gradient boosting scheme should be used.

## Adaboost (discrete version): assumptions

- binary classification task  $y \in \{+1, -1\}$
- family of base classifiers  $h(x) = h(x|\theta)$  where  $\theta$  is some fitted parametrization.
- $h(x) \in \{+1, -1\}$
- classification is performed with  
$$\hat{y} = sign\{f_0(x) + c_1 f_1(x) + \dots + c_M f_M(x)\}$$
- optimized loss is  $L(y, f(x)) = e^{-yf(x)}$
- FSAM is applied

## Adaboost (discrete version): algorithm

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; number of additive weak classifiers  $M$ , a family of weak classifiers  $h(x) \in \{+1, -1\}$ , trainable on weighted datasets.

- ➊ Initialize observation weights  $w_i = 1/n$ ,  $i = 1, 2, \dots, n$ .
  - ➋ for  $m = 1, 2, \dots, M$ :
    - ➌ fit  $h^m(x)$  to training data using weights  $w_i$
    - ➍ compute weighted misclassification rate:
$$E_m = \frac{\sum_{i=1}^N w_i \mathbb{I}[h^m(x_i) \neq y_i]}{\sum_{i=1}^N w_i}$$
  - ➎ if  $E_M > 0.5$  or  $E_M = 0$ : terminate procedure.
  - ➏ compute  $\alpha_m = \ln((1 - E_m)/E_m)$
  - ➐ increase all weights, where misclassification with  $h^m(x)$  was made:
- $$w_i \leftarrow w_i e^{\alpha_m}, i \in \{i : h^m(x_i) \neq y_i\}$$

**Output:** composite classifier  $f(x) = \text{sign} \left( \sum_{m=1}^M \alpha_m h^m(x) \right)$

# Table of Contents

## 1 Gradient boosting

# Motivation

- Problem: For general loss function  $L$  FSAM cannot be solved explicitly
- Analogy with function minimization: when we can't find optimum explicitly we use numerical methods
- Gradient boosting: numerical method for iterative loss minimization

# Gradient descent algorithm

$$F(w) \rightarrow \min_w, \quad w \in \mathbb{R}^N$$

Gradient descend algorithm:

**INPUT:**

$\eta$ -parameter, controlling the speed of convergence  
 $M$ -number of iterations

**ALGORITHM:**

initialize  $w$

**for**  $m = 1, 2, \dots M$ :

$$\Delta w = \frac{\partial F(w)}{\partial w}$$

$$w = w - \eta \Delta w$$

# Modified gradient descent algorithm

**INPUT:**

$M$ -number of iterations

**ALGORITHM:**

initialize  $w$

**for**  $m = 1, 2, \dots M$ :

$$\Delta w = \frac{\partial F(w)}{\partial w}$$

$$c^* = \arg \min_c F(w - c\Delta w)$$

$$w = w - c^* \Delta w$$

# Gradient boosting

- Now consider  $F(f(x_1), \dots, f(x_N)) = \sum_{n=1}^N \mathcal{L}(f(x_n), y_n)$
- Gradient descent performs pointwise optimization, but we need generalization, so we optimize in space of functions.
- Gradient boosting implements modified gradient descent in function space:
  - find  $z_i = -\frac{\partial \mathcal{L}(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - fit base learner  $h_m(x)$  to  $\{(x_i, z_i)\}_{i=1}^N$

## Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
  - ② fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
- ② fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$

- ③ solve univariate optimization problem:

$$\sum_{i=1}^N \mathcal{L}(f_{m-1}(x_i) + c_m h_m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
- ② fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$

- ③ solve univariate optimization problem:

$$\sum_{i=1}^N \mathcal{L}(f_{m-1}(x_i) + c_m h_m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

- ④ set  $f_m(x) = f_{m-1}(x) + c_m h_m(x)$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f_0(x)$  (might be taken  $f_0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$
- ② fit  $h_m$  to  $\{(x_i, z_i)\}_{i=1}^N$ , for example by solving

$$\sum_{n=1}^N (h_m(x_n) - z_n)^2 \rightarrow \min_{h_m}$$

- ③ solve univariate optimization problem:

$$\sum_{i=1}^N \mathcal{L}(f_{m-1}(x_i) + c_m h_m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

- ④ set  $f_m(x) = f_{m-1}(x) + c_m h_m(x)$

**Output:** approximation function  $f_M(x) = f_0(x) + \sum_{m=1}^M c_m h_m(x)$

## Gradient boosting: examples

In gradient boosting

$$\sum_{n=1}^N \left( h_m(x_n) - \left( -\frac{\partial \mathcal{L}(r, y_n)}{\partial r} \Big|_{r=f^{m-1}(x_n)} \right) \right)^2 \rightarrow \min_{h_m}$$

Commonly used loss-functions<sup>1</sup>:

- $\mathcal{L} = \frac{1}{2} (r - y)^2$
- $\mathcal{L} = e^{-ry}$
- $\mathcal{L} = [-ry]_+$
- $\mathcal{L} = \ln(1 + e^{-ry})$

---

<sup>1</sup>Estimate targets for them.

## Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N \mathcal{L}(\gamma, y_i)$$

## Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N \mathcal{L}(\gamma, y_i)$$

- ② For each step  $m = 1, 2, \dots, M$ :

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N \mathcal{L}(\gamma, y_i)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} |_{r=f^{m-1}(x_i)}$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N \mathcal{L}(\gamma, y_i)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$

- ② fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .

## Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N \mathcal{L}(\gamma, y_i)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$

- ② fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .

- ③ for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} \mathcal{L}(f_{m-1}(x_i) + \gamma, y_i)$$

## Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N \mathcal{L}(\gamma, y_i)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$

- ② fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .

- ③ for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} \mathcal{L}(f_{m-1}(x_i) + \gamma, y_i)$$

- ④ update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$

## Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, N$ ; loss function  $\mathcal{L}(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f_0(x)$ :

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N \mathcal{L}(\gamma, y_i)$$

- ② For each step  $m = 1, 2, \dots, M$ :

- ① calculate derivatives  $z_i = -\frac{\partial \mathcal{L}(r, y_i)}{\partial r} \Big|_{r=f^{m-1}(x_i)}$

- ② fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^N$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .

- ③ for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} \mathcal{L}(f_{m-1}(x_i) + \gamma, y_i)$$

- ④ update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$

**Output:** approximation function  $f_M(x)$

## Modification of boosting for trees

- Compared to first method of gradient boosting, boosting of regression trees finds additive coefficients individually for each terminal region  $R_{jm}$ , not globally for the whole classifier  $h^m(x)$ .
- This is done to increase accuracy: forward stagewise algorithm cannot be applied to find  $R_{jm}$ , but it can be applied to find  $\gamma_{jm}$ , because second task is solvable for arbitrary  $L$ .
- Max leaves  $J$ 
  - interaction between no more than  $J - 1$  terms
  - usually  $J \leq 8$
  - $M$  controls underfitting-overfitting tradeoff and selected using validation set

## Shrinkage & subsampling

- Shrinkage of general GB, step (d):

$$f_m(x) = f_{m-1}(x) + \nu c_m h_m(x)$$

- Shrinkage of trees GB, step (d):

$$f_m(x) = f_{m-1}(x) + \nu \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$$

- Comments:
  - $\nu \in (0, 1]$
  - $\nu \downarrow \implies M \uparrow$
- Subsampling
  - increases speed of fitting
  - may increase accuracy

# Types of boosting

- Loss function  $F$ :
  - $\mathcal{L}(|f(x) - y|)$  - regression
  - $\mathcal{L}(y \cdot \text{score}(y = +1|x))$  - binary classification
  - multiclass classification - may extend with one vs. all scheme
- Optimization
  - analytical (AdaBoost)
  - gradient based
- Base learners: continuous, discrete
- Extensions: shrinkage, subsampling